

01-Data Augmentation and Text Classification Foundation

1. A Survey on Data Augmentation for Text Classification — Bayer, Kaufhold and Reuter (2022)
Download/access: <https://arxiv.org/abs/2107.03158>
Publisher: <https://doi.org/10.1145/3544558>

02-Datasets

1. deepset Prompt-injections Dataset — deepset (2023)
Download/access: <https://huggingface.co/datasets/deepset/prompt-injections>

03-Core Transformer, BERT and Fine-Tuned Classification Models

1. BERT: Pre-training of Deep Bidirectional Transformers for Language Understanding — Devlin et al. (2019)
Download/access: <https://arxiv.org/abs/1810.04805>
ACL: <https://aclanthology.org/N19-1423/>
2. SecBERT: Privacy-preserving pre-training based neural network inference system — Huang and Wang (2024)
Download/access: <https://doi.org/10.1016/j.neunet.2024.106135>
PubMed: <https://pubmed.ncbi.nlm.nih.gov/38271920/>
3. RoBERTa: A Robustly Optimized BERT Pretraining Approach — Liu et al. (2019)
Download/access: <https://arxiv.org/abs/1907.11692>
4. DistilBERT, a distilled version of BERT: Smaller, faster, cheaper and lighter — Sanh et al. (2019)
Download/access: <https://arxiv.org/abs/1910.01108>
5. Prompt Injection Detection Using Ensemble of BERT and LSTM with GloVe Embeddings — Chisty et al. (2024)
Download/access: https://www.researchgate.net/publication/392638465_Prompt_Injection_Detection_Using_Ensemble_of_BERT_and_LSTM_with_GloVe_Embeddings
6. Securing IMF from Intent Injection Using BERT-Based Classifiers — Kim (2025)
Download/access: <https://doi.org/10.23919/ICMU65253.2025.11219161>
ICMU programme: <https://www.icmu.org/icmu2025/?program=>
7. Applying pre-trained multilingual BERT in embeddings for improved malicious prompt injection attacks detection — Rahman, Shahriar, Wu and Cuzzocrea (2024)
Download/access: <https://arxiv.org/abs/2409.13331>

04-Explainable AI and Model Interpretation

1. A Unified Approach to Interpreting Model Predictions — Lundberg and Lee (2017)
Download/access: <https://arxiv.org/abs/1705.07874>
Publisher:

<https://papers.nips.cc/paper/7062-a-unified-approach-to-interpreting-model-predictions>

2. Explainable Artificial Intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI — Arrieta et al. (2020)
Download/access: <https://arxiv.org/abs/1910.10045>
Publisher: <https://doi.org/10.1016/j.inffus.2019.12.012>
3. A Diagnostic Study of Explainability Techniques for Text Classification — Atanasova et al. (2020)
Download/access: <https://aclanthology.org/2020.emnlp-main.263/>
PDF: <https://aclanthology.org/2020.emnlp-main.263.pdf>
4. Enjoy the Saliency: Towards Better Transformer-Based Faithful Explanations with Word Saliency — Chrysostomou and Aletras (2021)
Download/access: <https://aclanthology.org/2021.emnlp-main.645/>
PDF: <https://aclanthology.org/2021.emnlp-main.645.pdf>
5. SHAP Values for Text Classification Tasks — CoderzColumn (2022)
Access:
<https://coderzcolumn.com/tutorials/artificial-intelligence/shap-values-for-text-classification-tasks>
6. ERASER: A Benchmark to Evaluate Rationalized NLP Models — DeYoung et al. (2020)
Download/access: <https://aclanthology.org/2020.acl-main.408/>
PDF: <https://aclanthology.org/2020.acl-main.408.pdf>
7. European Union regulations on algorithmic decision-making and a “right to explanation” — Goodman and Flaxman (2017)
Download/access: <https://arxiv.org/abs/1606.08813>
Publisher: <https://doi.org/10.1609/aimag.v38i3.2741>

05-LLM Foundation, Transformer Architecture and General LLM Background

1. Qwen Technical Report — Bai et al. (2023)
Download/access: <https://arxiv.org/abs/2309.16609>
2. Constitutional AI: Harmlessness from AI Feedback — Bai et al. (2022)
Download/access: <https://arxiv.org/abs/2212.08073>
3. Language Models are Few-Shot Learners — Brown et al. (2020)
Download/access: <https://arxiv.org/abs/2005.14165>
NeurIPS:
<https://proceedings.neurips.cc/paper/2020/hash/1457c0d6bfcb4967418bfb8ac142f64a-Abstract.html>
4. Unleashing the Potential of Prompt Engineering for Large Language Models — Chen, Zhang, Langrené and Zhu (2025)
Download/access: <https://arxiv.org/abs/2310.14735>
PubMed: <https://pubmed.ncbi.nlm.nih.gov/40575123/>
5. Switch Transformers: Scaling to Trillion Parameter Models with Simple and Efficient Sparsity — Fedus, Zoph and Shazeer (2022)

- Download/access: <https://jmlr.org/papers/v23/21-0998.html>
PDF: <https://jmlr.org/papers/volume23/21-0998/21-0998.pdf>
6. Mistral 7B — Jiang et al. (2023)
Download/access: <https://arxiv.org/abs/2310.06825>
 7. Retrieval-Augmented Generation for Knowledge-Intensive NLP Tasks — Lewis et al. (2020)
Download/access: <https://arxiv.org/abs/2005.11401>
NeurIPS PDF: https://proceedings.neurips.cc/paper_files/paper/2020/file/6b493230205f780e1bc26945df7481e5-Paper.pdf
 8. DeepSeek-V3 Technical Report — Liu et al. (2024)
Download/access: <https://arxiv.org/abs/2412.19437>
 9. Training Language Models to Follow Instructions with Human Feedback — Ouyang et al. (2022)
Download/access: <https://arxiv.org/abs/2203.02155>
OpenReview: <https://openreview.net/forum?id=TG8KACxEON>
 10. Improving Language Understanding by Generative Pre-Training — Radford et al. (2018)
Download/access: https://cdn.openai.com/research-covers/language-unsupervised/language_understanding_paper.pdf
 11. A Systematic Survey of Prompt Engineering in Large Language Models: Techniques and Applications — Sahoo et al. (2024)
Download/access: <https://arxiv.org/abs/2402.07927>
 12. Gemini: A Family of Highly Capable Multimodal Models — Gemini Team (2023)
Download/access: <https://arxiv.org/abs/2312.11805>
 13. Kimi k1.5: Scaling Reinforcement Learning with LLMs — Kimi Team (2025)
Download/access: <https://arxiv.org/abs/2501.12599>
 14. LLaMA: Open and Efficient Foundation Language Models — Touvron et al. (2023)
Download/access: <https://arxiv.org/abs/2302.13971>
 15. Attention is All You Need — Vaswani et al. (2017)
Download/access: <https://arxiv.org/abs/1706.03762>
NeurIPS: <https://papers.nips.cc/paper/7181-attention-is-all-you-need>
 16. Chain-of-Thought Prompting Elicits Reasoning in Large Language Models — Wei et al. (2022)
Download/access: <https://arxiv.org/abs/2201.11903>
NeurIPS: https://proceedings.neurips.cc/paper_files/paper/2022/hash/9d5609613524ecf4f15af0f7b31abca4-Abstract-Conference.html
 17. Qwen2 Technical Report — Yang et al. (2024)
Download/access: <https://arxiv.org/abs/2407.10671>

06-Prompt Injection, Jailbreak, Adversarial Prompting and LLM Security

1. Can Indirect Prompt Injection Attacks Be Detected and Removed? — Chen et al. (2025)
Download/access: <https://aclanthology.org/2025.acl-long.890/>
PDF: <https://aclanthology.org/2025.acl-long.890.pdf>
2. A Wolf in Sheep's Clothing: Generalized Nested Jailbreak Prompts Can Fool Large Language Models Easily — Ding et al. (2024)
Download/access: <https://aclanthology.org/2024.naacl-long.118/>
PDF: <https://aclanthology.org/2024.naacl-long.118.pdf>
3. Not What You've Signed Up For: Compromising Real-World LLM-Integrated Applications with Indirect Prompt Injection — Greshake et al. (2023)
Download/access: <https://arxiv.org/abs/2302.12173>
Publisher: <https://dl.acm.org/doi/10.1145/3605764.3623985>
4. DataSentinel: A Game-Theoretic Detection of Prompt Injection Attacks — Liu, Jia, Jia, Song and Gong (2025)
Download/access: <https://arxiv.org/abs/2504.11358>
IEEE: <https://www.computer.org/csdl/proceedings-article/sp/2025/223600c190/26EkG5ih2qQ>
5. FLIRT: Feedback Loop In-Context Red Teaming — Mehrabi et al. (2024)
Download/access: <https://aclanthology.org/2024.emnlp-main.41/>
PDF: <https://aclanthology.org/2024.emnlp-main.41.pdf>
6. LLM Prompt Injection AML.T0051 — MITRE ATLAS (2026)
Access: <https://atlas.mitre.org/techniques/AML.T0051>
7. OWASP Top 10 for LLM Applications — OWASP (2025)
Access: <https://genai.owasp.org/llm-top-10/>
Prompt Injection page: <https://genai.owasp.org/llmrisk/llm01-prompt-injection/>
8. Neural Exec: Learning and Learning from Execution Triggers for Prompt Injection Attacks — Pasquini, Strohmeier and Troncoso (2024)
Download/access: <https://arxiv.org/abs/2403.03792>
Publisher: <https://dl.acm.org/doi/10.1145/3689932.3694764>
Code: https://github.com/pasquini-dario/LLM_NeuralExec
9. Automatic Prompt Optimization with Gradient Descent and Beam Search — Pryzant et al. (2023)
Download/access: <https://aclanthology.org/2023.emnlp-main.494/>
PDF: <https://aclanthology.org/2023.emnlp-main.494.pdf>
10. Ignore This Title and HackAPrompt: Exposing Systemic Vulnerabilities of LLMs through a Global Prompt Hacking Competition — Schulhoff et al. (2023)
Download/access: <https://aclanthology.org/2023.emnlp-main.302/>
PDF: <https://aclanthology.org/2023.emnlp-main.302.pdf>
11. Do Anything Now: Characterizing and Evaluating In-the-Wild Jailbreak Prompts on Large Language Models — Shen et al. (2024)
Download/access: <https://arxiv.org/abs/2308.03825>
12. LinkPrompt: Natural and Universal Adversarial Attacks on Prompt-Based Language Models — Xu and Wang (2024)

- Download/access: <https://aclanthology.org/2024.naacl-long.360/>
PDF: <https://aclanthology.org/2024.naacl-long.360.pdf>
13. A Novel Evaluation Framework for Assessing Resilience Against Prompt Injection Attacks in Large Language Models — Yip, Esmradi and Chan (2023)
Download/access: <https://arxiv.org/abs/2401.00991>
DOI: <https://doi.org/10.1109/CSDE59766.2023.10487667>
14. How Johnny Can Persuade LLMs to Jailbreak Them — Zeng et al. (2024)
Download/access: <https://aclanthology.org/2024.acl-long.773/>
PDF: <https://aclanthology.org/2024.acl-long.773.pdf>
15. Adaptive Attacks Break Defenses Against Indirect Prompt Injection Attacks on LLM Agents — Zhan et al. (2025)
Download/access: <https://aclanthology.org/2025.findings-naacl.395/>
PDF: <https://aclanthology.org/2025.findings-naacl.395.pdf>
16. InjecAgent: Benchmarking Indirect Prompt Injections in Tool-Integrated Large Language Model Agents — Zhan et al. (2024)
Download/access: <https://aclanthology.org/2024.findings-acl.624/>
PDF: <https://aclanthology.org/2024.findings-acl.624.pdf>
17. PromptRobust: Towards Evaluating the Robustness of Large Language Models on Adversarial Prompts — Zhu et al. (2023)
Download/access: <https://arxiv.org/abs/2306.04528>
Publisher: <https://dl.acm.org/doi/10.1145/3689217.3690621>
18. PromptGuard: A Structured Framework for Injection-Resilient Language Models — Alzahrani (2026)
Download/access: <https://www.nature.com/articles/s41598-025-31086-y>
19. Defending Against Alignment-Breaking Attacks via Robustly Aligned LLM — Cao et al. (2024)
Download/access: <https://aclanthology.org/2024.acl-long.568/>
PDF: <https://aclanthology.org/2024.acl-long.568.pdf>
20. A Hitchhiker's Guide to Jailbreaking ChatGPT via Prompt Engineering — Liu et al. (2024)
Download/access: <https://dl.acm.org/doi/10.1145/3663530.3665021>
Preprint/PDF: <https://tianweiz07.github.io/Papers/24-SEA4DQ.pdf>
21. Enhancing System Security: LLM-Driven Defense Against Prompt Injection Vulnerabilities — Muliarevych (2024)
Download/access: https://www.researchgate.net/publication/386127152_Enhancing_System_Security_LLM-Driven_Defense_Against_Prompt_Injection_Vulnerabilities
22. Prompt Injection Attack Detection with Machine Learning — Özçam et al. (2025)
Access: <https://www.semanticscholar.org/paper/Prompt-Injection-Attack-Detection-with-Machine-%C3%96z%C3%A7am-Solmaz/9851a84b4749556c02f677a5a9f7c8c7a09c39ea>

23. Prompt Analysis in Large Language Models: Adversarial Prompt Detection — Sriram et al. (2026)
Access: <https://rdems.in/index.php/RDEMS/article/view/126>
24. A Real-Time Framework for Prompt Injection Attacks Detection on Cloud-Hosted Large Language Models — Thaqi, Martiri and Rexha (2025)
Access: <https://doi.org/10.1109/FLLM67465.2025.11391009>
25. Preemptive Answer Attacks on Chain-of-Thought Reasoning — Xu, Qi and Xu (2024)
Download/access: <https://aclanthology.org/2024.findings-acl.876/>
PDF: <https://aclanthology.org/2024.findings-acl.876.pdf>

07-Prompt, Vision-Language and Red-Teaming Related Extra Papers

1. One Prompt Word is Enough to Boost Adversarial Robustness for Pre-Trained Vision-Language Models — Li et al. (2024)
Download/access: <https://arxiv.org/abs/2403.01849>
Code: <https://github.com/TreeLLi/APT>
2. Adversarial Nibbler: An Open Red-Teaming Method for Identifying Diverse Harms in Text-to-Image Generation — Quaye et al. (2024)
Download/access: <https://arxiv.org/abs/2403.12075>
Publisher: <https://dl.acm.org/doi/10.1145/3630106.3658913>

08-Market, Statistics and Industry Context

1. Enterprise LLM Market Size and Share, Statistics Report 2025–2034 — Global Market Insights (2025)
Access: <https://www.gminsights.com/industry-analysis/enterprise-llm-market>
2. Prompt Injection Statistics 2026: Hidden Risks Now — SQ Magazine (2026)
Access: <https://sqmagazine.co.uk/prompt-injection-statistics/>

09-Medical, Psychology, Education and Domain-Specific Extra Papers

1. Large Language Models in Medical Chatbots: Opportunities, Challenges, and the Need to Address AI Risks — Chow and Li (2025)
Download/access: <https://www.mdpi.com/2078-2489/16/7/549>
DOI: <https://doi.org/10.3390/info16070549>
2. PsycOLLM: Enhancing LLM for Psychological Understanding and Evaluation — Hu et al. (2024)
Download/access: <https://arxiv.org/abs/2407.05721>
Project: <https://github.com/MindIntLab-HFUT/PsycOLLM>
3. Evaluation of AI-Based Chatbots in Liver Cancer Information Dissemination — Karaagac and Carkit (2025)
Download/access: <https://doi.org/10.1159/000546726>
Publisher:

<https://karger.com/ocl/article/104/3/283/928648/Evaluation-of-AI-Based-Chatbots-in-Liver-Cancer>

4. **A Systematic Review of AI, VR, and LLM Applications in Special Education — Voultsiou and Moussiades (2025)**

Download/access: <https://link.springer.com/article/10.1007/s10639-025-13550-4>